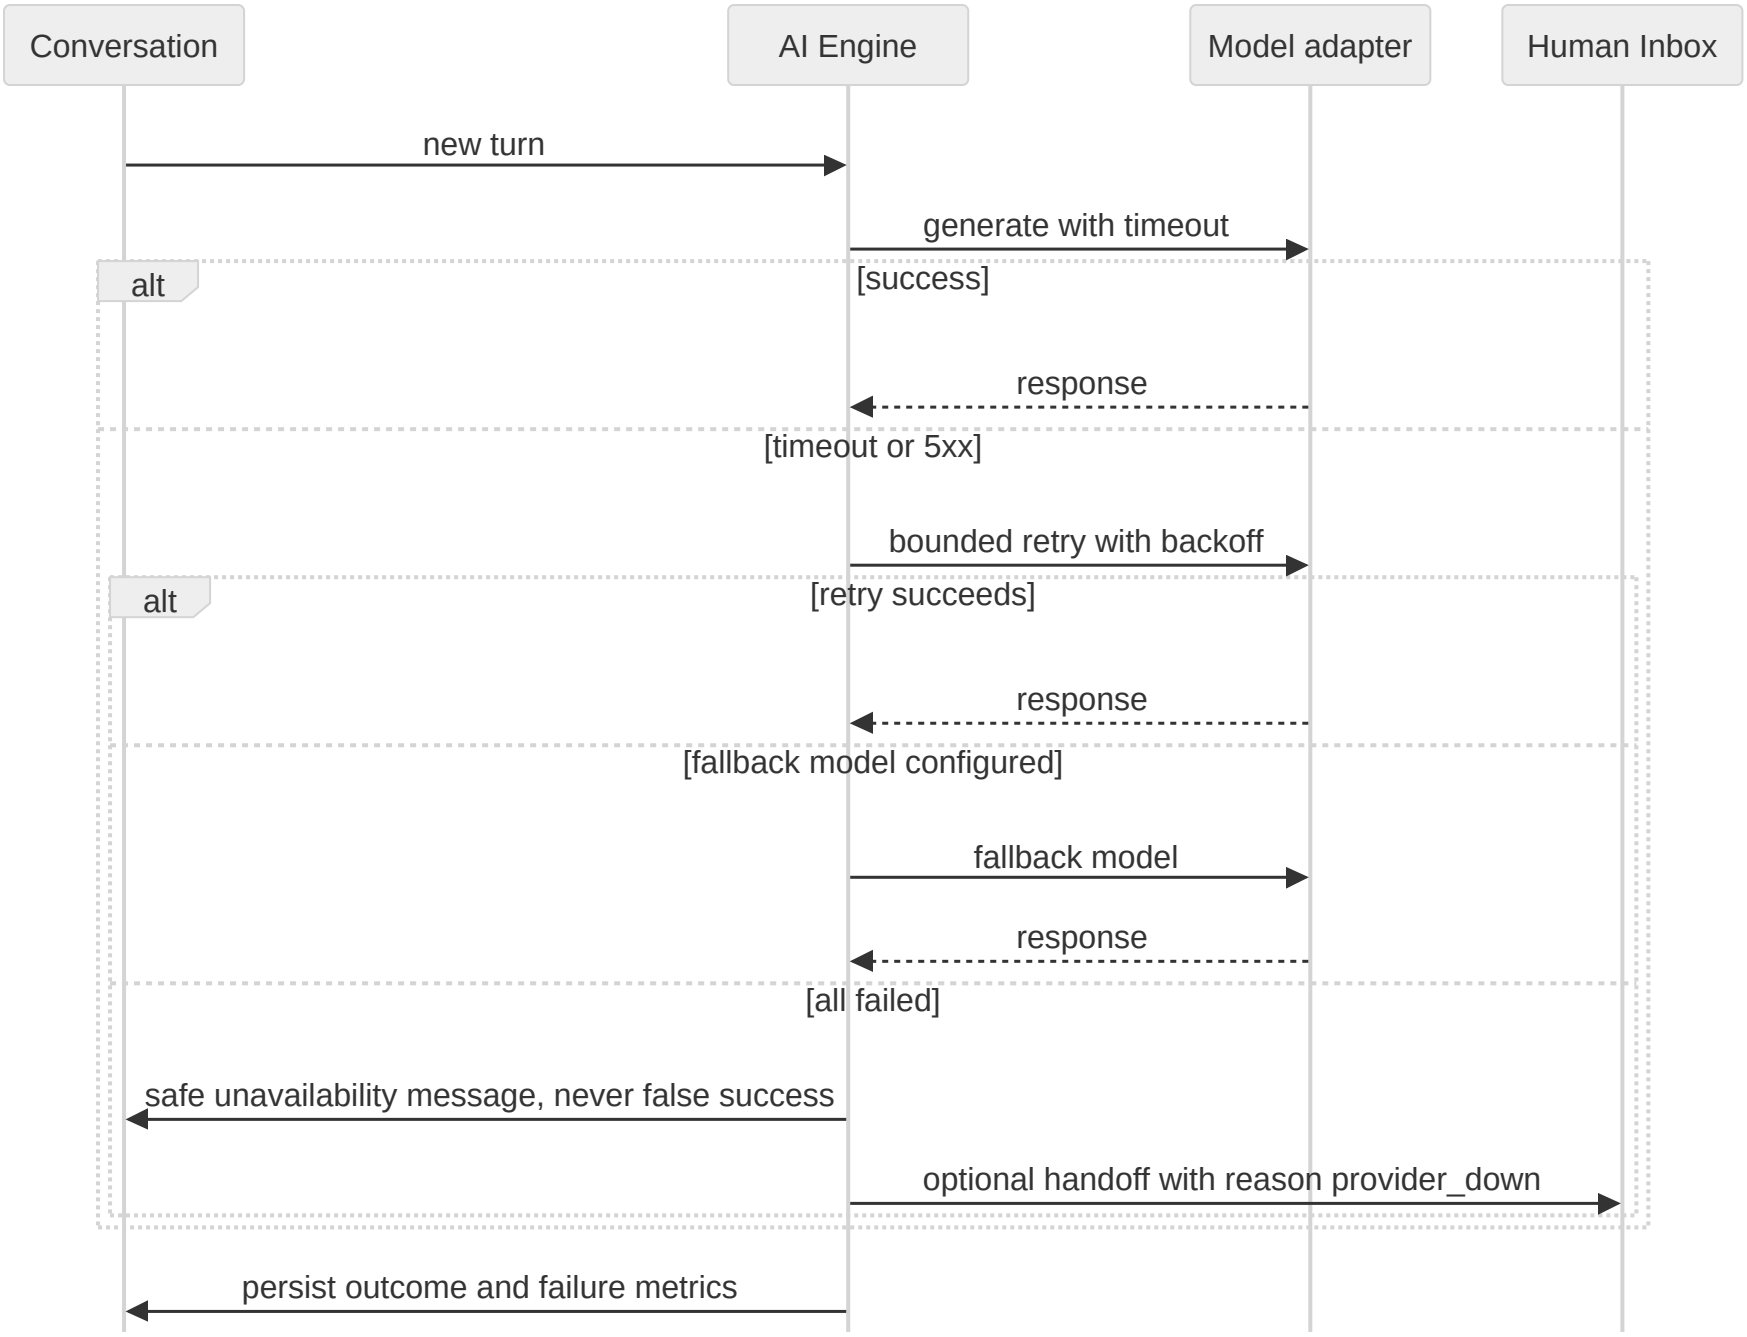
Conversation
AI Engine
Model adapter
Human Inbox
new turn
generate with timeout
alt
[success]
response
[timeout or 5xx]
bounded retry with backoff
alt
[retry succeeds]
response
[fallback model configured]
fallback model
response
[all failed]
safe unavailability message, never false success
optional handoff with reason provider_down
persist outcome and failure metrics