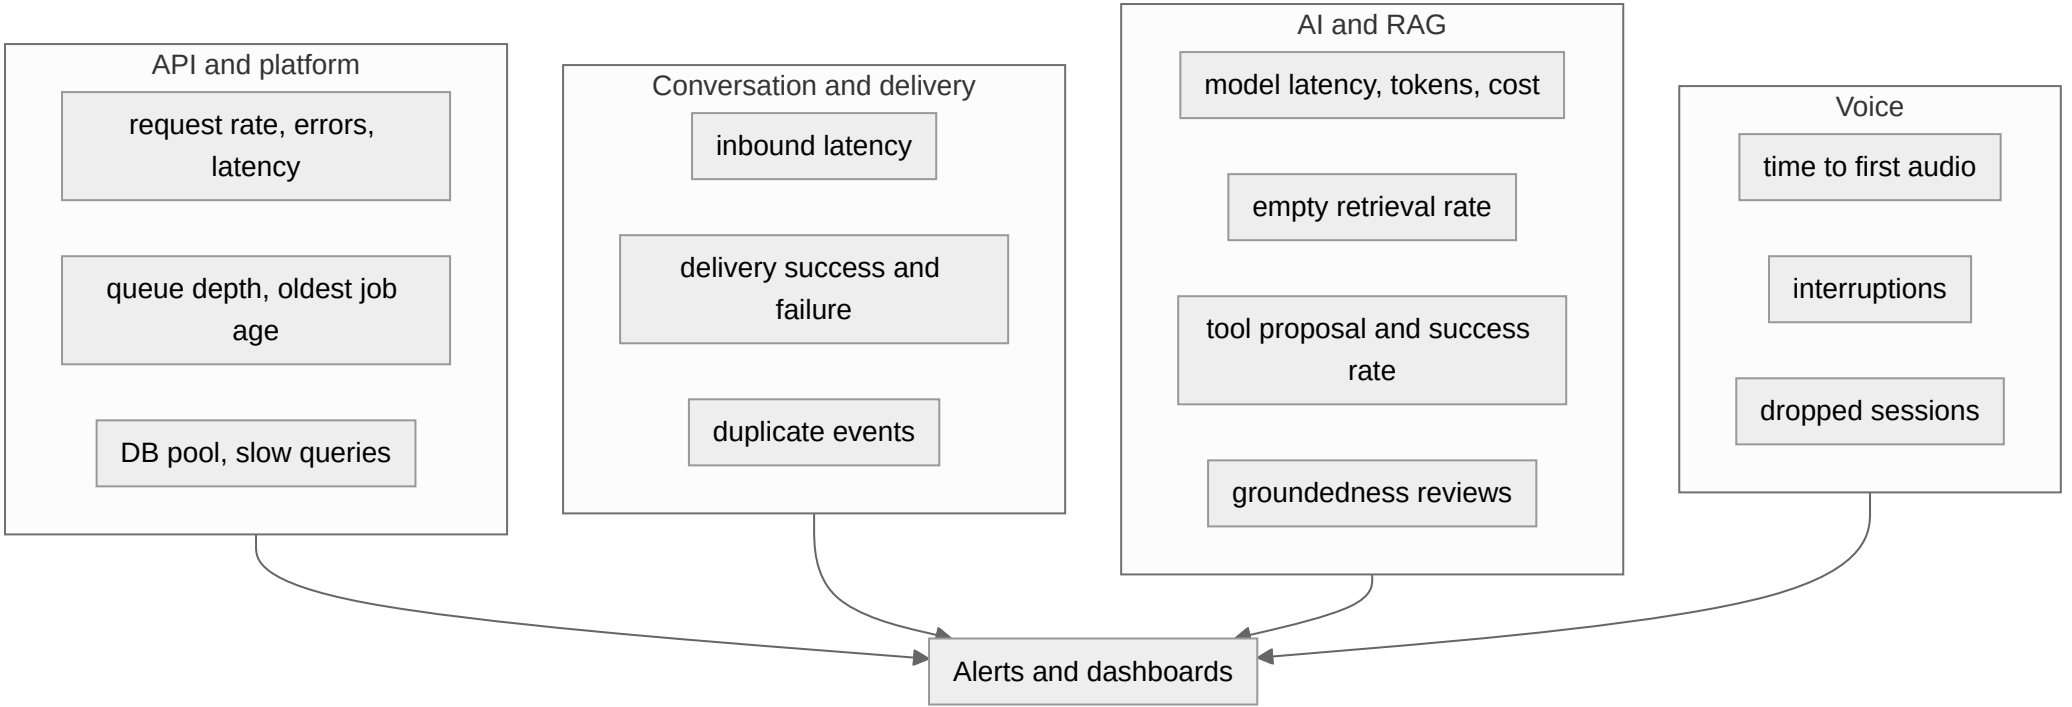
API and platform
request rate, errors, latency
queue depth, oldest job age
DB pool, slow queries
Conversation and delivery
inbound latency
delivery success and failure
duplicate events
AI and RAG
model latency, tokens, cost
empty retrieval rate
tool proposal and success rate
groundedness reviews
Voice
time to first audio
interruptions
dropped sessions
Alerts and dashboards