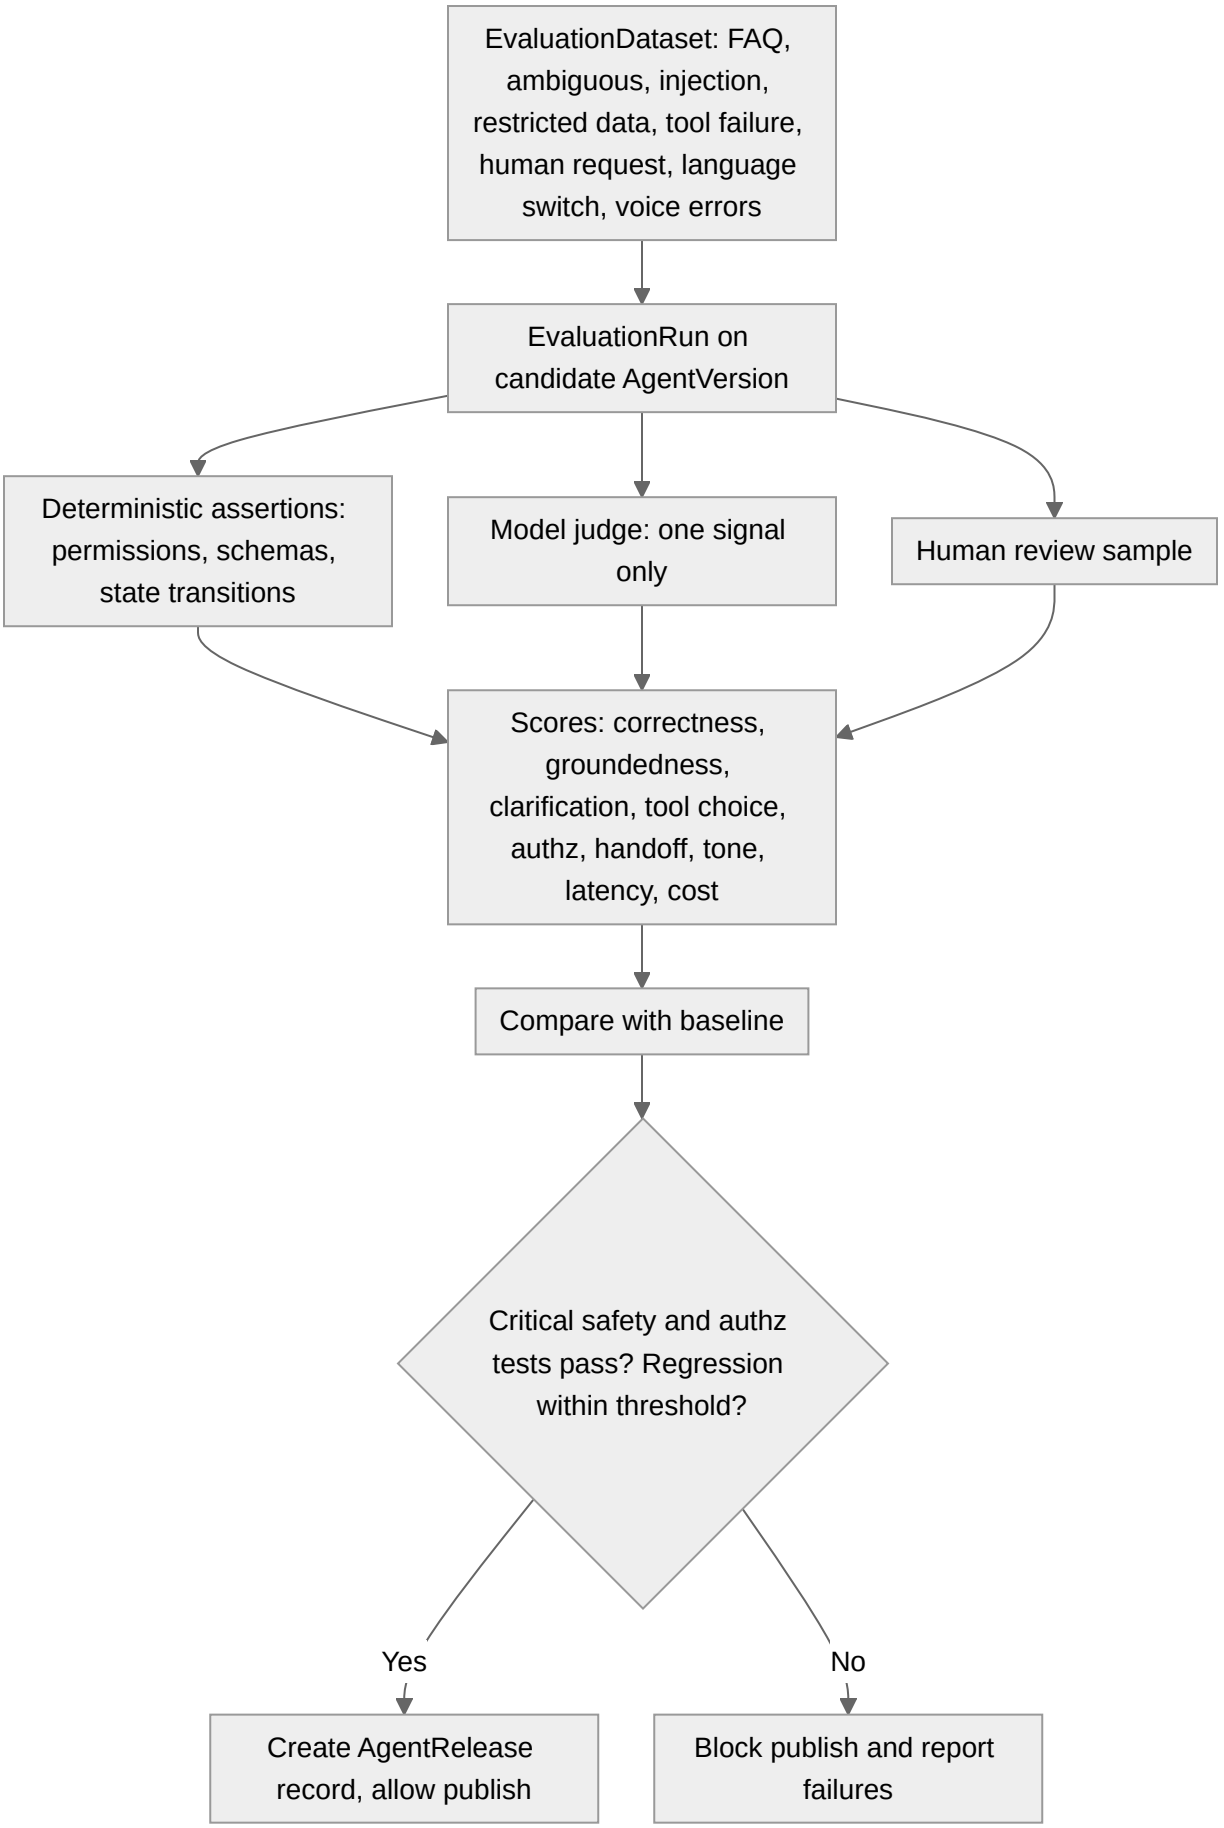
EvaluationDataset: FAQ, ambiguous, injection, restricted data, tool failure, human request, language switch, voice errors
EvaluationRun on candidate AgentVersion
Deterministic assertions: permissions, schemas, state transitions
Model judge: one signal only
Human review sample
Scores: correctness, groundedness, clarification, tool choice, authz, handoff, tone, latency, cost
Compare with baseline
Critical safety and authz tests pass? Regression within threshold?
Yes
No
Create AgentRelease record, allow publish
Block publish and report failures