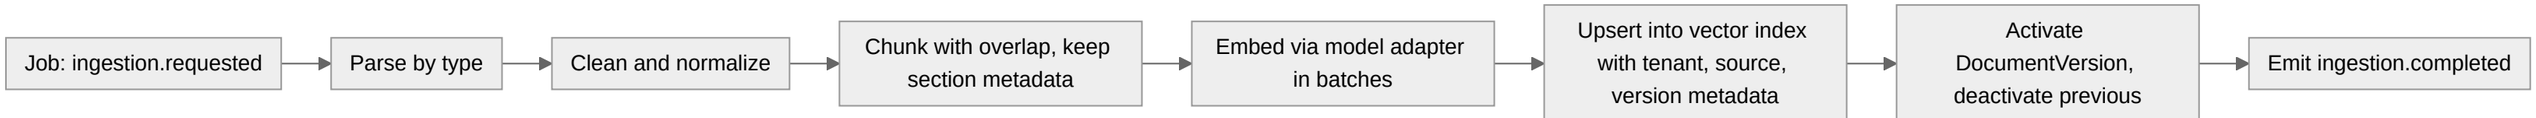
Job: ingestion.requested
Parse by type
Clean and normalize
Chunk with overlap, keep section metadata
Embed via model adapter in batches
Upsert into vector index with tenant, source, version metadata
Activate DocumentVersion, deactivate previous
Emit ingestion.completed